

# Efficient Parallel C++ A practical course

Tobias Maier, Sascha Witt  
Winter 2021/22

INSTITUTE OF THEORETICAL INFORMATICS · ALGORITHMICS GROUP



```
179 .L9:
180     vmovdqa ymm1, YMMWORD
181     add rdx, 1
182     vpmullw ymm0, ymm1, ymm1
183     vpmulhw ymm1, ymm0, ymm1
184     vpunpcklwd ymm2, ymm0, ymm1
185     vpunpckhwd ymm1, ymm2, ymm1
186     vperm2i128 ymm0, ymm2, ymm1, 32
187     vperm2i128 ymm2, ymm0, ymm1, 49
188     vpand ymm0, ymm3, ymm2
189     vpand ymm2, ymm0, ymm0
190     vpackusdw ymm0, ymm0, 216
191     vpermq ymm0, ymm0, YMMWORD PTR
192     vmovdqu ymm0, rdx
193     add rcx, 32
194     cmp r10, rdx
195     ja .L9
196     mov rdx, rsi
197     and rdx, -16
198     lea r8, [rdx+rdx]
199     mov rax, [r9+r8]
```

## After this course, you should...

- be able to write efficient, parallel, clean, maintainable, testable C++
- know what and how to optimize
  - Testing and profiling
  - Cache efficiency, branch mispredictions, data structures,...
- be aware of common performance pitfalls
  - Copies, indirection, virtual functions,...
- know about common pitfalls of parallelization
  - False sharing, excessive synchronization,...

Cheat sheet available (we'll come back to that)

## There will be 5 exercises:

- |   |             |
|---|-------------|
| 1. Common bottlenecks, templates                  | 2 weeks     |
| 2. Cache efficiency, testing, profiling           | 2 weeks     |
| 3. Parallelization using OpenMP                   | 2 weeks     |
| 4. Parallelization using <code>std::thread</code> | 3 weeks     |
| 5. Personal project                               | 1 + 4 weeks |

- Processing times overlap by half a week.
- Exact topics are subject to change.
- You will need  $\frac{1}{3}$  of the points of *each* exercise to pass.
- Turn-in consists of code, experimental evaluation and a short interview where we ask about your work.

# Schedule

## November

1	2	3	4	5	6	7
8	9	10	11	12	13	14
15	16	17	18	19	20	21
22	23	24	25	26	27	28
29	30					

## December

		1	2	3	4	5
6	7	8	9	10	11	12
13	14	15	16	17	18	19
20	21	22	23	24	25	26
27	28	29	30	31		

## January

					1	2
3	4	5	6	7	8	9
10	11	12	13	14	15	16
17	18	19	20	21	22	23
24	25	26	27	28	29	30
31						

## February

	1	2	3	4	5	6
7	8	9	10	11	12	13
14	15	16	17	18	19	20
21	22	23	24	25	26	27
28						

```
https://git.scc.kit.edu/ITI10/efficient-parallel-cpp
```

## We use Git

- to distribute the exercises
- for turn-in
  - Fork the repo, then give *us* access to *your* repo.
  - Do not check in large files (i.e. graph data)!

## Performance graphs

- R (e.g., *R for Data Science* [<http://r4ds.had.co.nz/>])
- Python (e.g., *Matplotlib* [<https://matplotlib.org/>])
- Gnuplot
- ...

## Code quality

- clang-format
- clang-tidy
- ...

# The most important rules

## Always measure.

- It's not an optimization if it doesn't make your code faster.

## Focus.

- In inner loops, a single instruction can make a huge difference.
- A function that accounts for only 3% of your running time cannot give you more than a 3% improvement.

## Know what the compiler can and cannot do for you.

- And then verify that it is actually doing what you expect it to do.
- Compiler Explorer [<https://godbolt.org>] can be helpful.

## Unnecessary copies

- Difference between *object*, *reference*, and *pointer*
- Call by *value*, *reference*, or *pointer*
- Pass small types by value.
- Prefer passing large types by *(const)* reference.
- move semantics are another way to reduce unnecessary copies.

## Indirection

- Indirection often leads to cache misses.
- cf. `std::list` **VS** `std::vector`

## Heap allocation

- It's slow.
- It adds indirection.
- cf. `std::array` **VS** `std::vector`

## Virtual functions

- vtables are a level of indirection.
- Virtual function calls generally cannot be inlined.
  - The `final` specifier can help with this.
- Templates can provide similar functionality, at compile time.



## Cache efficiency

- Memory access is often a bottleneck.
- Things that can improve locality of reference include:
  - Less indirection
  - Linear (or predictable) access patterns
  - Different data representation (e.g., Structure of Arrays)

## Branch mispredictions

- Modern CPU pipelines are long, and clearing them is expensive.
- Things that can help the branch predictor include:
  - Avoiding (data-dependent) branches
  - Vectorization

## Data races

- Just don't.
- Use atomics, or locks where that is not possible.

## False sharing

- Sometimes, not sharing is caring.
- one cache line has 64byte this is the minimum offset between two objects to avoid false sharing  
(`std::hardware_destructive_interference_size`).

## Unnecessary synchronization

- Minimize shared data and communication.
- Work on local data, then combine in a separate step.

# Problem-specific parallelization issues

## Use local approximations instead of global data.

- Think about what you need to know locally about global data.

## Distribute data efficiently.

- This is especially important on Non-Uniform Memory Access (NUMA) architectures.

## Distribute work efficiently.

- Equal amounts of data do not imply equal amounts of work.
- Some workloads can be highly asymmetrical.