

Efficient Parallel C++ A practical course

Tobias Maier, Sascha Witt
Winter 2021/22

INSTITUTE OF THEORETICAL INFORMATICS · ALGORITHMICS GROUP



```
179 .L9: vmovdqa ymm1, YMMWORD
180 add rdx, 1
181 vpmullw ymm0, ymm1, ymm1
182 vpmulhw ymm1, ymm0, ymm1
183 vpunpcklwd ymm1, ymm0, ymm1
184 vpunpckhwd ymm0, ymm2, ymm1
185 vperm2i128 ymm2, ymm0, ymm1, 32
186 vperm2i128 ymm2, ymm2, ymm0
187 vpand ymm0, ymm3, ymm2
188 vpand ymm2, ymm0, ymm0
189 vpackusdw ymm0, ymm0, 216
190 vpermq ymm0, ymm0, YMMWORD PTR [r8+rcx], ymm0
191 vmovdqu ymm1, YMMWORD PTR [r8+rcx], ymm0
192 add rcx, 32
193 cmp r10, rdx
194 ja .L9
195 mov rdx, rsi
196 and rdx, -16
197 lea r8, [rdx+rdx]
```

Advantages

- Widely supported standard
- Portable
- Easy to use

Disadvantages

- Fixed model of execution (fork-join)
- Some loss of control

Parallel region

```
#pragma omp parallel
{
    // parallel code
}
```

Parallel loop

```
#pragma omp parallel for
for (int i = 0; i < n; ++i) {
    // parallel loop
}
```

Shared and private variables

```
#pragma omp parallel private(x, y) shared(z)
{
    // parallel code
}
```

std::thread

Advantages

- Standard since C++11
- Portable
- Control over thread lifetimes

Disadvantages

- Less out-of-the-box functionality
- Easier to get wrong

Single thread

```
std::thread t{[] { std::cout << "Hello,_world!\n"; }};  
// do other stuff  
t.join();
```

Multiple threads

```
std::vector<std::thread> threads;  
for (int i = 0; i < num_threads; ++i)  
    threads.emplace_back([i] {  
        std::cout << "Hello_from_thread_" << i << "!\n";  
    });  
// do other stuff  
for (auto& t : threads)  
    t.join();
```

`std::async`

- Forget about it until you *really* understand what it does.

No standard support yet, need to roll your own

```
class Worker {  
    public:  
        Worker() : thread_([this] {  
            while (!stopped()) {  
                // wait for a new task to execute and run it  
            }  
        }) {}  
  
        // ...  
  
    private:  
        std::thread thread_;  
        std::function<void()> task_;  
        // ...  
};
```

Never simply access data that may be concurrently modified.

- All data is shared by all threads by default.
- You can use `thread_local` to make a private copy for each thread.

Use synchronization when accessing data that may be modified.

- Prefer atomics if possible.
- Use locks when you need to modify larger chunks of data.
- You also need to use synchronization when *reading*.

Required to avoid data races under concurrent writes

```
std::atomic<int> x(0);  
x += 2;  
// _almost_ the same as  
x.fetch_add(2);
```

Compare-and-swap operations

```
std::atomic<int> x(0);  
  
int old_value = x;  
do {  
    int new_value = do_something_with(old_value);  
} while (!x.compare_exchange_weak(old_value, new_value));
```


Heavy-duty locking

```
std::mutex mtx;  
  
// ...  
  
{    // acquire exclusive access  
    std::lock_guard lock(mtx);  
  
    // do stuff  
  
}    // auto-release at end of scope
```

Use `std::lock` to avoid deadlocks

```
std::lock(mtx3, mtx1, mtx2);  
  
// do stuff  
  
// don't forget about unlocking!  
// or use std::lock_guard with std::adopt_lock
```

Waiting for other threads

Spinlocks: Quick-and-dirty busy waiting

```
std::atomic_flag flag = ATOMIC_FLAG_INIT;  
// ...  
  
while (flag.test_and_set(std::memory_order_acquire)) {}  
// do stuff  
flag.clear(std::memory_order_release);
```

Condition Variables

```
std::condition_variable cv;  
  
std::unique_lock lk(mtx);  
cv.wait(lk, [&] { return ready; });  
// ...  
  
{  
    std::unique_lock lk(mtx);  
    ready = true;  
}  
cv.notify_one();
```

Some caches are shared among cores and sockets

- L1 and L2 is usually private to each core (but shared by hyper-cores!).
- L3 is usually shared by all cores on a socket.
- RAM is obviously shared by all sockets.

Shared caches can be good

- If multiple cores work with the same data, only one has to pay the cost of loading it into the cache.
- Only really works if you don't modify anything.

Cache misses are expensive

- If you change a cache line, everybody else has to throw it away.
- That means they have to reload it before they can access it again.

Put unrelated, frequently-updated data on different cache lines

```
struct alignas(64) SharedData {  
    // ...  
};  
std::vector<SharedData> vec;
```

...or even on different pairs of cache lines

```
struct alignas(128) SharedData {  
    // ...  
};
```

Non-Uniform Memory Access

- Only appears in multi-socket machines.
- Each socket has fast access to some part of the main memory.
- When accessing other parts, latency and/or bandwidth is worse.
- There are libraries that help manipulating where memory is located.
- You probably don't need to worry too much about it for this course.

Hyper-Threading

- Commonly supported by all modern processors.
- One core may “simultaneously” execute two threads.
- Execution is interleaved to fill dead time (due to, e.g., memory latency).
- Usually leads to some speedup, but sometimes hurts overall performance.

Limiting which cores can execute certain threads

- Usually, the operating system can move your threads around as it wishes.
- Sometimes, this can lead to bad performance.
- Especially on NUMA systems, pinning threads to sockets or even specific cores can be a good idea.
- Can also lead to worse performance if not done carefully.

When writing parallel code:

- **Always** be thread-safe: Use atomics, or locks if necessary.
- Use spinlocks if the expected waiting time is low, otherwise use mutexes (even better, but more complicated: spin for a while, then switch to a mutex).
- Think about cache effects (especially false sharing).
- Multi-socket scenarios are complicated.